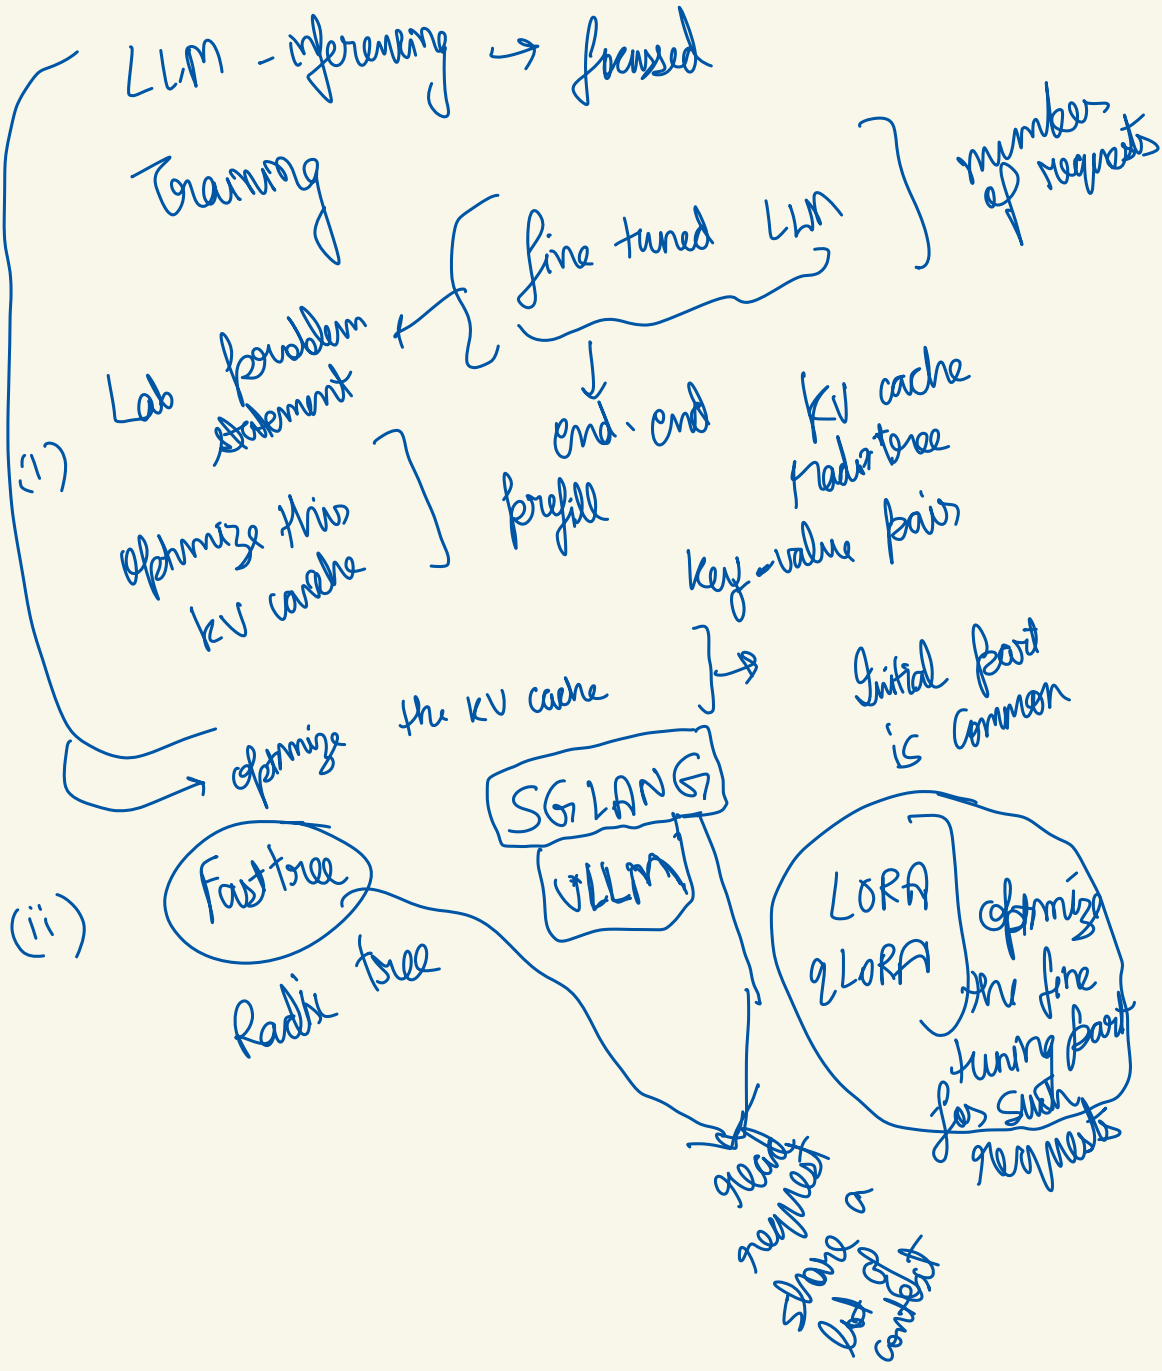

LLM - inferencing → focussed
Training
fine tuned LLM
number of requests
Lab problem statement
optimize this kv cache
(i)
end-end prefill
KV cache radix tree
key-value pairs
optimize
the KV cache
Initial part is common
SGLANG
VLLM
(ii)
Fast tree
Radix tree
LORA
QLORA
optimize the fine tuning part for such requests
group requests sharing a lot of context